

Every answer cited to an approved source, *or refused*.

TrustOps answers security questionnaires from your own approved evidence. Deterministic gates decide what ships, not model confidence: every released answer carries a citation, every refusal names its gap and routes to the right human, and nothing reaches the delivered workbook without a named reviewer in a tamper evident audit log.

0

UNSUPPORTED CLAIMS, ENFORCED
RELEASE RULE

~2 min

YOUR DOCUMENTS TO TEN GATED
ANSWERS

Public

ADVERSARIAL EVAL SUITE, ON
GITHUB

01 THE FOUR PLANTED TRAPS, REFUSED CORRECTLY ON EVERY RUN

■ CERT NEVER INFERRED

Certification asked, roadmap on file. Plans never qualify as certificates, however confident the draft sounds.

■ SOURCES CONFLICT

Two approved policies disagree. Both quarantined until an owner reconciles; questions citing either refuse too.

■ EVIDENCE EXPIRED

Stale pentest report. Cannot support a current claim; the refusal names the document, expiry, and owner.

■ ROUTED TO COUNSEL

Unlimited liability demanded. Legal commitments route before drafting; the model never sees them.

02 HOW IT WORKS

01

Upload

Policies and reports. You attest approval.

02

Retrieve

Your documents only. Tenants isolated.

03

Draft

A model drafts. The gates never trust it.

04

Gate

Cite or refuse. Stale, conflict, legal.

05

Review

A named human releases, on the chain.

WHAT LEAVES YOUR INFRASTRUCTURE

Documents stay in an isolated workspace and are deleted after 14 days, never used for training. Retrieval runs locally on our server. The drafting model sees one question and its excerpts, never your full documents, never another tenant.

WHAT YOU TAKE AWAY

Your workbook returned in its original structure, an audit working paper, and a hash chained log you can verify yourself.

```
DELIVERED.xlsx / run_report.html /  
audit_log.jsonl
```

Run it on your own documents. No call, no form.

github.com/dhruvshahi10/trustops
dhruv.shahi07@gmail.com